

Reverse Item Response Theory for Sparsity-Robust Ranking in Fragmented Cancer Drug-Response Matrices

Jung Min Kang*
Independent Researcher
Seoul, South Korea

May 2026

Abstract

We introduce reverse Item Response Theory (IRT) to pharmacogenomic drug-response analysis by treating cancer types as latent “subjects” with resistance ability and drugs as “items” with evasion difficulty. Applied to 242,036 drug sensitivity measurements from the Genomics of Drug Sensitivity in Cancer (GDSC2) database, the model estimates cancer-type-level in-vitro resistance and drug-level broad activity on a shared latent scale. Validation across four missingness regimes demonstrates that reverse IRT better recovers the full-data latent ranking than simple averaging, with advantages of $\Delta\rho = +0.089$ to $+0.095$ at 60% missingness under MCAR, cancer-biased, and drug-biased sparsity. Held-out prediction confirms IRT achieves the best Brier score among five evaluated methods. Bootstrap confidence intervals show 19 of 28 cancer types have stable resistant/sensitive classifications. Cross-platform PRISM replication shows a weak, non-significant rank-order correlation ($\rho = 0.25$, $p = 0.33$, 17 cancer types), indicating the contribution is methodological robustness under fragmented evaluation, not a universal clinical resistance leaderboard.

1 Introduction

Large-scale pharmacogenomic screening efforts, including the Genomics of Drug Sensitivity in Cancer [GDSC; Yang et al., 2012, Iorio et al., 2016] and the PRISM Repurposing dataset [Corsello et al., 2020], have generated comprehensive drug-response matrices spanning hundreds of drugs and thousands of cancer cell lines. Existing approaches include ANOVA-based biomarker discovery [Garnett et al., 2012], machine learning prediction [Costello et al., 2014], deep learning [Liu et al., 2020], and recent work inferring general principles of drug sensitivity with experimental validation [Carli et al., 2025].

These approaches model each drug–cell-line pair independently or predict sensitivity from genomic features, without jointly estimating a cancer type’s global resistance and a drug’s global activity on a common measurement scale.

Item Response Theory [IRT; Lord and Novick, 1968, Baker and Kim, 2004] provides this capability. IRT jointly estimates subject ability and item difficulty on a shared latent scale. Kang [2026a] demonstrated that IRT-based ranking outperforms simple averaging under sparse evaluation in AI benchmarking. Rodriguez et al. [2021] applied IRT to NLP evaluation, and Polo et al. [2024] used IRT for efficient LLM benchmarking.

*Email: gangjeongmin23@gmail.com

We propose a conceptual inversion: cancer types become “subjects” with latent resistance θ_j , drugs become “items” with evasion difficulty b_i :

$$P(\text{sensitive}) = \sigma(b_i - \theta_j) \quad (1)$$

The primary contribution is not the specific rankings—which are platform-dependent—but demonstrating that reverse IRT provides sparsity-robust ranking recovery in fragmented drug-response matrices. This is relevant because real-world therapeutic evidence matrices are sparse: drugs, indications, and trial populations are unevenly evaluated.

2 Data

GDSC2. Release 8.5 (October 2023), Wellcome Sanger Institute. Raw: 242,036 drug–cell-line measurements (969 cell lines, 286 drugs, 32 TCGA cancer types). After removing unclassified types (196,345 remaining) and excluding CLL (9 drugs, insufficient coverage), we aggregate to a 28×286 matrix with 7,821 cells (97.7% coverage). Binarization: sensitive if fitted $\text{LN_IC50} < 3.297$ (global median, $\approx 27 \mu\text{M}$). A global threshold avoids drug-specific normalization; drug parameters reflect apparent broad in-vitro activity rather than absolute pharmacological potency.

PRISM. Secondary dose-response dataset [Corsello et al., 2020]: 701,004 IC50 entries, 1,448 compounds, 499 cell lines. Cell lines mapped via DepMap lineage metadata, yielding 17 overlapping cancer types.

3 Methods

Reverse IRT. 1PL model (Eq. 1) with analytical gradients (finite-difference error < 0.02), L-BFGS-B optimization, Gaussian priors ($\sigma^2 = 4$). Parameters: $28 + 286 = 314$ for 7,821 observations.

Sparsity test. Remove 20–60% of cells under four regimes: MCAR, cancer-biased (harder cancers lose more), drug-biased (weaker drugs lose more), pathway-block (entire pathways removed). Evaluate Spearman ρ vs. full-data IRT ranking (15 seeds).

Held-out prediction. 10 independent random hold-out splits, each holding out 20% of observed cells (repeated random splits, not k -fold cross-validation). Brier score against cancer-only averaging, drug-only averaging, two-way additive (global + cancer effect + drug effect), logistic fixed effects, and reverse IRT.

Bootstrap CIs. 200 drug-panel (column) resamples for 95% CIs on θ .

PRISM replication. Independent reverse IRT on PRISM; Spearman ρ with GDSC2 ranking.

4 Results

4.1 Sparsity Robustness

Reverse IRT outperforms averaging across all regimes (Table 1, Figure 1). The advantage generally increases with missingness, except under pathway-block missingness where it remains positive but modest.

4.2 Held-Out Prediction

IRT achieves the best Brier score among all five evaluated methods (Table 2, Figure 2).

Table 1: $\Delta\rho = \text{IRT} - \text{averaging}$ (difference of the 15-seed mean ρ) for recovery of the full-data IRT ranking under induced sparsity. $\pm$: SD of the IRT ρ across the 15 seeds (per-seed $\Delta\rho$ values were not saved).

Regime	20%	40%	60%
MCAR	$+0.028 \pm 0.003$	$+0.049 \pm 0.003$	$+0.089 \pm 0.008$
Cancer-biased	$+0.025 \pm 0.002$	$+0.048 \pm 0.006$	$+0.095 \pm 0.016$
Drug-biased	$+0.033 \pm 0.002$	$+0.066 \pm 0.005$	$+0.093 \pm 0.006$
Pathway-block	$+0.008 \pm 0.004$	$+0.019 \pm 0.004$	$+0.018 \pm 0.008$




Figure 1: Ranking recovery advantage ($\Delta\rho$) of reverse IRT over averaging under four missingness regimes. Error bars show ± 1 SD of the IRT ρ across 15 seeds (not the SD of $\Delta\rho$). IRT outperforms averaging in all conditions.

4.3 Cancer Resistance Ranking

The ranking shows face-valid concordance with known clinical difficulty patterns (Figure 3). Pancreatic adenocarcinoma (PAAD) ranks most resistant, directionally consistent with clinical difficulty (13.7% five-year relative survival; SEER). Hematological malignancies rank most sensitive, consistent with therapeutic advances in ALL ($\sim 90\%$ childhood cure rate; NCI PDQ). Nineteen of 28 cancer types have stable resistant/sensitive classifications (CIs not crossing zero); nine middle-tier cancers remain uncertain.

4.4 External Replication

PRISM replication yields a weak, non-significant rank-order correlation, $\rho = 0.252$ ($p = 0.33$; 17 overlapping cancer types; Figure 4). We do not use directional (sign) agreement as evidence: every PRISM θ is positive (minimum 0.274), so sign agreement only counts the cancers with GDSC2 $\theta > 0$ (14/17) and carries no information from PRISM. The three sign mismatches are simply

Table 2: Brier score over held-out cells from 10 independent random 20% hold-out splits (lower = better).

Method	Brier	Δ vs IRT
Cancer-only averaging	0.1338	+0.1195
Drug-only averaging	0.0340	+0.0197
Two-way additive	0.0176	+0.0032
Logistic fixed effects	0.0181	+0.0038
Reverse IRT	0.0143	—




Figure 2: Held-out Brier score among five evaluated methods (10 random 20% hold-out splits). Reverse IRT outperforms all methods including the fairer two-way additive and logistic fixed-effect baselines.

the three overlapping cancers with GDSC2 $\theta < 0$ (STAD, HNSC, NB). STAD and HNSC have bootstrap CIs crossing zero; NB does not (95% CI $[-0.914, -0.277]$), so NB is not in the bootstrap uncertainty zone.




Figure 3: Cancer in-vitro resistance with bootstrap 95% CIs (200 drug-panel resamples). Red: highly resistant ($\theta > 0.5$); orange: moderately resistant; light blue: moderately sensitive; dark blue: highly sensitive. CIs crossing zero indicate uncertain classification.




Figure 4: GDSC2 vs PRISM resistance estimates (17 overlapping cancer types). Blue: GDSC2 $\theta > 0$; red crosses: GDSC2 $\theta < 0$. All PRISM θ are positive, so sign agreement is uninformative and is not reported. Only extremes and the GDSC2 $\theta < 0$ cancers are labeled; full table in repository outputs. Rank-order correlation is weak and non-significant ($\rho = 0.25$, $p = 0.33$).

5 Discussion

The primary finding is methodological: reverse IRT provides sparsity-robust ranking recovery in drug-response matrices. The Evaluation Failure Scaling Law mechanism [Kang, 2026a], originally demonstrated in AI benchmark evaluation, transfers to pharmacogenomic data.

Limitations. (1) Cell-line in-vitro resistance does not equal clinical resistance. (2) Binarization at a global LN₅₀ median discards continuous information. (3) The 1PL model assumes unidimensional resistance; the LLTM [Fischer, 1973, Kang, 2026b] with mutation features could decompose resistance into interpretable components. (4) PRISM replication is weak and non-significant in rank order ($\rho = 0.25$, $p = 0.33$), possibly reflecting platform differences.

6 Conclusion

Reverse IRT provides a sparsity-robust framework for ranking cancer types and drugs on a shared latent scale. The method outperforms averaging under all tested missingness regimes, achieves the best held-out calibration among five evaluated methods, and produces rankings with face-valid clinical concordance. The contribution is methodological: when drug-response matrices become fragmented, reverse IRT preserves ranking structure better than averaging. Cross-platform replication confirms this is a ranking methodology contribution, not a universal biological discovery.

Data and Code Availability

GDSC2 Release 8.5 is available from CancerRxGene:

https://www.cancerrxgene.org/downloads/bulk_download

PRISM secondary dose-response and DepMap cell line metadata are available from DepMap:

<https://depmap.org/repurposing>

All code, validation scripts, and output CSVs:

<https://github.com/testofschool/reverse-irt-cancer>

Running `challenge_c_v3_final.py` and `validation_suite.py` also writes the run-metadata files `metadata.json` and `validation_metadata.json`; these files are not included in the repository.

References

- Baker, F. B. and Kim, S.-H. (2004). *Item Response Theory*. Marcel Dekker.
- Carli, F. et al. (2025). Learning and actioning general principles of cancer cell drug sensitivity. *Nat. Commun.*, 16, 1654.
- Corsello, S. M. et al. (2020). Discovering the anticancer potential of non-oncology drugs. *Nat. Cancer*, 1, 235–248.
- Costello, J. C. et al. (2014). A community effort to assess and improve drug sensitivity prediction. *Nat. Biotechnol.*, 32, 1202–1212.
- Fischer, G. H. (1973). The linear logistic test model. *Acta Psychol.*, 37, 359–374.

- Garnett, M. J. et al. (2012). Systematic identification of genomic markers of drug sensitivity. *Nature*, 483, 570–575.
- Iorio, F. et al. (2016). A landscape of pharmacogenomic interactions in cancer. *Cell*, 166, 740–754.
- Kang, J. M. (2026a). The scaling law of evaluation failure. *arXiv:2605.11205*.
- Kang, J. M. (2026b). Explaining benchmark difficulty: LLTM for feature-based AI evaluation. Preprint.
- Liu, Q. et al. (2020). DeepCDR: hybrid graph convolutional network for cancer drug response. *Bioinformatics*, 36, i911–i918.
- Lord, F. M. and Novick, M. R. (1968). *Statistical Theories of Mental Test Scores*. Addison-Wesley.
- Polo, F. M. et al. (2024). Efficient multi-prompt evaluation of LLMs. *NeurIPS 2024*.
- Rodriguez, P. et al. (2021). Evaluation examples are not equally informative. *ACL-IJCNLP*, 4486–4503.
- Yang, W. et al. (2012). Genomics of Drug Sensitivity in Cancer. *Nucleic Acids Res.*, 41, D955–D961.